




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Advanced C++ for HPC: STL

Anton Afanasyev, CSCS (afanasyev@cscs.ch)

2-4 October, 2017

Resources:

Definition

- **STL** stands for **S**tandard **T**emplate **L**ibrary
- Originally was designed by Alexander Stepanov
- The spec and implementation went to public in 1994
- Became one of the key stones of the C++ Standard Library
- Consists of containers, algorithms, iterators, functors and related stuff

Design principles

- Concepts serve as generalized interfaces
 - Each STL (template) class implements a concept or several concepts
 - For each STL template class parameter it is defined what concept(s) it should satisfy
- Orthogonal and extensible
 - Iterators are concepts
 - Containers implement iterators
 - Algorithms operate on iterators
 - Applying algorithms to containers (or parts of them) creates combinatorial number of usage possibilities
- No runtime overhead
- Everything is thread compatible but not thread safe

Iterator Concept

Iterators should behave as pointers within C-array

```
int data[] = { 1, 2, 3, 4 };  
for (auto i = data; i != data + 4; ++i)  
    std::cout << *i << "\n";
```

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

++i, i++	move forward
--i, i--	move backward
i+5, i-5, i-j, i+=5, i-=5	move anywhere
*i	dereference
i == j, i != j	compare
i = j	(cheap) copy

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector<int> data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

`++i, i++` move forward

`--i, i--` move backward

`i+5, i-5, i-j, i+=5, i-=5` move anywhere

`*i` dereference

`i == j, i != j` compare

`i = j` (cheap) copy

RandomAccessIterator

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector<int> data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

<code>++i, i++</code>	move forward
<code>--i, i--</code>	move backward
<code>i+5, i-5, i-j, i+=5, i-=5</code>	move anywhere
<code>*i</code>	dereference
<code>i == j, i != j</code>	compare
<code>i = j</code>	(cheap) copy

BidirectionalIterator

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector<int> data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

<code>++i, i++</code>	move forward
<code>--i, i--</code>	move backward
<code>i+5, i-5, i-j, i+=5, i-=5</code>	move anywhere
<code>*i</code>	dereference
<code>i == j, i != j</code>	compare
<code>i = j</code>	(cheap) copy

ForwardIterator

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector<int> data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

<code>++i, i++</code>	move forward, but only once
<code>--i, i--</code>	move backward
<code>i+5, i-5, i-j, i+=5, i-=5</code>	move anywhere
<code>*i</code>	dereference, but only for reading
<code>i == j, i != j</code>	compare
<code>i = j</code>	(cheap) copy

InputIterator

Iterator Concepts

Iterators should behave as pointers within C-array

```
std::vector<int> data = { 1, 2, 3, 4 };  
for (auto i = data.begin(); i != data.end(); ++i)  
    std::cout << *i << "\n";
```

<code>++i, i++</code>	move forward, but only once
<code>--i, i--</code>	move backward
<code>i+5, i-5, i-j, i+=5, i-=5</code>	move anywhere
<code>*i</code>	dereference, but only for writing
<code>i == j, i != j</code>	compare
<code>i = j</code>	(cheap) copy

OutputIterator

Functors – yet another concept

Functors should behave like pointers to functions

$f(42)$		callable
$g = f$		(cheap) copy
		an invocation doesn't change the functor

Container Concept (simplified)

Container owns a sequence of values and provides access to them via iterators

`c.begin()`



`c.end()`

`c.size()`, `c.empty()` etc

Container Concept (possible alternative)

Container owns a sequence of values and provides access to them via iterators

`begin(c)`

value	value	value	value	value	value	value	value	value	value
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

`end(c)`

`size(c), empty(c)` etc

SequenceContainers

<code>std::array<T, N></code>	static contiguous array
<code>std::vector<T></code>	dynamic contiguous array
<code>std::deque<T></code>	double-ended queue
<code>std::forward_list<T></code>	singly-linked list
<code>std::list<T></code>	doubly-linked list

std::array

is like a C-array but better

```
// can be returned from a function
std::array<int, 42> foo();
```

```
// can act as a tuple
using obj_t = std::array<int, 3>;
obj_t obj = { 3, 4, 5 };
auto second = std::get<1>(obj);
static_assert(std::tuple_size<obj_t>{} == 3, "");
static_assert(std::is_same<std::tuple_element<0, obj_t>::type, int>{}, "");
```

```
// is 'true' sequence container
std::sort(obj.begin(), obj.end());
```


std::vector

Consists of a pointer to the buffer, logical size and capacity of the buffer



When capacity is not enough to fit newly inserted elements, a new buffer is allocated with doubled capacity; the data is moved into the new buffer and the old data is released.

```
std::vector<int> a;  
for (int i = 0; i != N; ++i) a.push_back(i);
```

How many times each element is moved on average?

vector

Consists of a pointer to the buffer, logical size and capacity of the buffer



When capacity is not enough to fit newly inserted elements, a new buffer is allocated with doubled capacity; the data is moved into the new buffer and the old data is released.

```
std::vector<int> a;  
for (int i = 0; i != N; ++i) a.push_back(i);
```

How many times each element is moved on average? **2**

std::vector: Raw Access

```
size_t get_results_size();  
void get_results(size_t size, struct Result* buf);  
  
std::vector<Result> wrapper() {  
    std::vector<Result> res(get_results_size());  
    get_results(res.size(), res.data());  
    return res;  
}
```

`std::vector<bool>`

possibly space-efficient specialization

- elements are not stored in contiguous buffer
- not thread compatible
- was added to the library as a proof of concept for template specialization and proxy pattern


avoid using it without a solid reason

std::deque

data is in chunks; chunks are referenced from TOC

- effective push_front/push_back without invalidating refs and iterators
- random access requires two derefs


shift = 3
chunk_size = 6



std::list

doubly-linked list; size is cached

- fast insertions/deletions, has splice()
- no random access
- refs and iterators remains valid even when splicing between list
- some algorithms are implemented as methods



Sequence containers: what to choose

- use `std::vector` as a default – it is the fastest and consumes less memory
- use `std::deque` if you are implementing some sort of queue (`pop_front/push_front`)
- use `std::list` if performance doesn't matter much but your design requires references and iterator to be always valid.

AssociativeContainer

- `std::set<T>`
- `std::map<K, V>`
- `std::multiset<T>`
- `std::multimap<K, V>`
- `std::unordered_set<T>`
- `std::unordered_map<K, V>`
- `std::unordered_multiset<T>`
- `std::unordered_multimap<K, V>`

AssociativeContainer

- `std::set<T, Comp>`
- `std::map<K, V, Comp>`
- `std::multiset<T, Comp>`
- `std::multimap<K, V, Comp>`
- `std::unordered_set<T>`
- `std::unordered_map<K, V>`
- `std::unordered_multiset<T>`
- `std::unordered_multimap<K, V>`

AssociativeContainer

- `std::set<T, Comp>`
- `std::map<K, V, Comp>`
- `std::multiset<T, Comp>`
- `std::multimap<K, V, Comp>`
- `std::unordered_set<T, Hash, Eq>`
- `std::unordered_map<K, V, Hash, Eq>`
- `std::unordered_multiset<T, Hash, Eq>`
- `std::unordered_multimap<K, V, Hash, Eq>`

AssociativeContainer requirements

- if you use non-default `Compare` functor, verify that it is sane
- values in `std::set` should be virtually immutable

insert VS operator[] VS at

```
std::map<std::string, int> m = {{"a", 1}, {"b", 2}, {"c", 3}};  
m.insert({"a", 11});  
m.at("b") = 22;  
m["c"] = 33;  
m["d"] = 44;  
// m.at("e") = 55;  
for(auto && v : m)  
    std::cout << v.first << ":" << v.second << "\n";
```

insert VS operator[] VS at

```
std::map<std::string, int> m = {{"a", 1}, {"b", 2}, {"c", 3}};  
m.insert({"a", 11});  
m.at("b") = 22;  
m["c"] = 33;  
m["d"] = 44;  
// m.at("e") = 55;  
for(auto && v : m)  
    std::cout << v.first << ":" << v.second << "\n";
```

```
a:1  
b:22  
c:33  
d:44
```

insert VS operator[] VS at

```
std::map<std::string, int> m = {{"a", 1}, {"b", 2}, {"c", 3}};  
m.insert({"a", 11});  
m.at("b") = 22;  
m["c"] = 33;  
m["d"] = 44;  
m.at("e") = 55;  
for(auto && v : m)  
    std::cout << v.first << ":" << v.second << "\n";
```

std::out_of_range: map::at: key not found

Algorithms

- `std::sort`
- `std::stable_sort`
- `std::transform`
- `std::merge`
- `std::partition`
- `std::accumulate`
- `std::iota`
- `std::remove`
- ...

using std::remove

```
std::vector v = {1, 2, 3, 4};  
std::remove(v.begin(), v.end(), 2);  
assert(v.size() == 3);
```


using std::remove

```
std::vector v = {1, 2, 3, 4};  
v.erase(std::remove(v.begin(), v.end(), 2), v.end());  
assert(v.size() == 3);
```

How to use `std::swap`?

```
struct A {  
    bool val;  
    friend void swap(A&, A&) { std::cout << "I refuse to swap!\n"; }  
};  
  
int main() {  
    bool x = false, y = true;  
    std::swap(x, y);  
    assert(x);  
    A a{false}, b{true};  
    std::swap(a, b);  
    assert(!a.val);  
}
```

How to use `std::swap`?

```
struct A {  
    bool val;  
    friend void swap(A&, A&) { std::cout << "I refuse to swap!\n"; }  
};  
  
int main() {  
    using std::swap;  
    bool x = false, y = true;  
    swap(x, y);  
    assert(x);  
    A a{false}, b{true};  
    swap(a, b);  
    assert(!a.val);  
}
```

`std::queue`, `std::stack`, `std::priorityqueue`

- they are container adaptors (use other containers underneath)
- `std::queue`: FIFO, push into one end and pop from the other
- `std::stack`: FILO, push and pop into/from one end
- `std::priority_queue`
 - ordered container, get the largest value in constant time
 - slow insertion/removal (logarithmic)

What is default container for those adaptors?

Why `pop()` returns void?

New Features: `std::initializer_list` Support

```
std::vector<int> eval(std::vector<int> src) {  
    // do smth with src.  
    return src;  
}  
  
int main() {  
    // automatic conversion to the right type  
    std::vector<double> a = {1., true, 3};  
    // assign operator also works with IL  
    a = {};  
    // IL can contain any expressions. The execution order is defined.  
    int i = 0;  
    std::vector v = {++i, ++i, ++i};  
    for(auto x : v) std::cout << x << "\n";  
    // nested IL for nested containers work  
    std::set<std::vector<int>> b = {  
        {},  
        {1, 2},  
        {3}  
    };  
    // IL ctors are implicit. If function param is container, IL can be used as an args.  
    auto c = eval({1, 2, 3});  
}
```

New Features: Move Semantics Support

- Container element type requirements are soften from copyable to movable
- If the type is not even movable, it is still possible to store it in STL containers by wrapping into `std::unique_ptr`.
- There is no more reason to store raw pointers in containers anymore.

New Features: Range Based Loop Syntax

- this language feature is tailored to use with STL by design
- it implicitly defines a new concept
- makes simple algorithms like *std::for_each* and *std::transform* useless

New Features: Perfect Forwarding

all insertion methods now have emplace variations

```
struct A {  
    A(int a, int b, int c) {}  
};
```


```
int main() {  
    std::vector<A> v;  
    v.push_back(A{1, 2, 3});  
    v.emplace_back(1, 2, 3);  
}
```




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Thank you for your attention.